

## ✓ Assignment Deep Learning Frameworks (Module 2)

### Theoretical Questions

#### Q1. What is Tensor flow 2.0, and how is different from Tensorflow 1.x?

Ans-> TensorFlow 2.0 is a major update to the open-source machine learning library TensorFlow, focusing on ease of use, improved productivity, and better integration with Python. It was released by Google to address some of the complexities and pain points experienced with TensorFlow 1.x.

- Key Differences from TensorFlow 1.x:
- Eager Execution by Default:
- TensorFlow 1.x: Required users to define a static computational graph first and then execute it within a `tf.Session`. This "define and run" paradigm could be less intuitive and harder to debug.
- TensorFlow 2.0: Enables eager execution by default, meaning operations are executed immediately and return concrete values, similar to how NumPy operates. This makes debugging easier and provides a more Pythonic development experience.
- API Cleanup and Simplification:
- TensorFlow 1.x: Contained many redundant or overlapping APIs, leading to confusion. Global variables and implicit namespaces were also prevalent.
- TensorFlow 2.0: Features a cleaner and more consistent API. Many redundant APIs were removed or consolidated, and Keras is now the primary high-level API for building models, promoting a more streamlined workflow.
- No More Sessions and Placeholders (mostly):
- TensorFlow 1.x: Relied heavily on `tf.Session` for graph execution and `tf.placeholder` for feeding data into the graph.
- TensorFlow 2.0: With eager execution, sessions and placeholders are largely eliminated for typical workflows, simplifying data input and model execution.
- `tf.function` for Graph Compilation:
- TensorFlow 1.x: Required explicit graph construction for performance optimization.
- TensorFlow 2.0: Introduces `tf.function` decorator, which automatically converts Python functions into TensorFlow graphs for performance benefits, while maintaining the eager execution feel during development.
- Improved Variable Management:
- TensorFlow 1.x: Variable management could be cumbersome, often requiring explicit variable scopes.
- TensorFlow 2.0: Keras layers and models provide convenient properties to manage variables locally within their scope, making variable tracking and management more intuitive.
- In essence, TensorFlow 2.0 aims to make the framework more accessible and user-friendly, particularly for beginners, while retaining the power and scalability for advanced applications through features like `tf.function`.

#### Q2. How do you install Tensorflow 2.0?

Ans-> TensorFlow 2.0 can be installed using pip, Python's package installer.

- Steps for Installation: Ensure Python and pip are installed: TensorFlow 2.0 requires Python 3.4 or later. Pip is typically included with Python installations. Open your terminal or command prompt. Install TensorFlow 2.0:
- For the CPU version: Code `pip install tensorflow==2.0.0`

#### Q3. What is the primary function of the `tf.function` in Tensorflow 2.0?

Ans-> The primary function of `tf.function` in TensorFlow 2.0 is to convert regular Python functions into callable TensorFlow graphs, thereby enabling performance optimization and portability.

- Specifically, `tf.function` achieves this by:
- Graph Compilation: It traces the Python function's execution with sample inputs to build a static TensorFlow graph. This graph represents the computation as a series of operations and data dependencies, allowing for various optimizations not possible with pure eager execution.
- Performance Enhancement: By compiling the function into a graph, TensorFlow can apply optimizations such as operation fusion, memory allocation optimization, and efficient execution on various hardware (CPUs, GPUs, TPUs). This leads to significantly faster execution, especially for repetitive computations within training loops or inference.
- Portability and Deployment: The generated graph is a self-contained representation of the computation, making it independent of the Python environment. This allows for easier deployment of models to different platforms and environments, including TensorFlow Lite for mobile/edge devices and TensorFlow Serving for production deployments.
- Integration with AutoGraph: `tf.function` seamlessly integrates with AutoGraph, which automatically transforms Python control flow statements (like `if`, `for`, `while`) into equivalent TensorFlow graph operations, enabling dynamic behavior within the compiled graph.

#### Q4. What is the purpose of the model class in Tensorflow 2.0?

Ans-> In TensorFlow 2.0, the `tf.keras.Model` class serves as the central abstraction for building and managing machine learning models. Its primary purposes are:

- **Encapsulation of Layers and Functionality:** A `Model` object groups together a network of layers, defining the overall architecture and how data flows through it. It encapsulates the forward pass (how input data is transformed into output predictions) and can also include custom training logic.
- **Built-in Training and Evaluation Loops:** The `Model` class provides convenient methods like `fit()`, `evaluate()`, and `predict()` for handling the standard machine learning workflow of training, evaluating performance, and making predictions on new data. This simplifies the process for users.
- **Handling Complex Architectures:** Unlike the simpler `Sequential` model, `Model` (especially when used with the Functional API or subclassing) allows for the creation of complex network \*\*\* topologies, including: Multiple inputs and outputs. Shared layers across different parts of the network. Non-linear connections like skip connections or residual connections.
- **Serialization and Saving:** `Model` objects can be easily saved and loaded, including their architecture, weights, and optimizer state. This enables users to persist trained models for later use or deployment without needing to retrain them.
- **Extensibility and Customization:** While providing built-in functionalities, the `Model` class is also highly extensible. Users can subclass `tf.keras.Model` to define custom training steps, loss calculations, or metrics, offering fine-grained control over the training process.

## Q5. How do you create a neural network using Tensorflow 2.0?

Ans-> Creating a neural network using TensorFlow 2.x primarily involves using the Keras API, which is integrated directly into TensorFlow. The process typically follows these steps: Import TensorFlow and Keras. Python

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
```

**Prepare your Data:** Load and preprocess your dataset. This often involves tasks like normalization (e.g., scaling pixel values to 0-1 range for images) and splitting into training and testing sets. **Define the Model Architecture:** Sequential API: For simple, layer-by-layer models, use `keras.Sequential`. Python

```
model = keras.Sequential([
    layers.Dense(64, activation='relu', input_shape=(input_dim,)),
    layers.Dense(32, activation='relu'),
    layers.Dense(num_classes, activation='softmax')
])
```

**Functional API:** For more complex models with multiple inputs/outputs or shared layers, use the Functional API. Python

```
input_tensor = keras.Input(shape=(input_dim,))
x = layers.Dense(64, activation='relu')(input_tensor)
output_tensor = layers.Dense(num_classes, activation='softmax')(x)
model = keras.Model(inputs=input_tensor, outputs=output_tensor)
```

**Compile the Model:** Configure the learning process by specifying the optimizer, loss function, and metrics. Python

```
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy', # or 'categorical_crossentropy' for one-hot
              metrics=['accuracy'])
```

**Train the Model:** Fit the model to your training data. Python

```
model.fit(x_train, y_train, epochs=10, batch_size=32, validation_data=(x_val, y_val))
```

**Evaluate the Model:** Assess the model's performance on unseen data. Python

```
test_loss, test_acc = model.evaluate(x_test, y_test, verbose=2)
print(f'\nTest accuracy: {test_acc}')
```

**Make Predictions:** Use the trained model to predict on new data. Python

```
predictions = model.predict(new_data)
```

## Q6. What is the importance of Tensor Space in Tensorflow?

Ans-> In TensorFlow, "Tensor Space" refers to the conceptual environment or framework where tensors exist and are manipulated. The importance of Tensor Space in TensorFlow lies in its fundamental role as the core data structure and operational unit for all computations within the framework.

- Here's why Tensor Space is crucial:
- **Fundamental Data Representation:** Tensors are the primary way data is represented in TensorFlow. They are multi-dimensional arrays that can hold various types of data, from scalar values (rank 0 tensors) to vectors, matrices, and higher-dimensional arrays, making them suitable for representing complex datasets like images, audio, and text.
- **Enabling Efficient Computation:** TensorFlow is designed to perform operations on tensors efficiently, especially on specialized hardware like GPUs and TPUs. The tensor-based architecture allows for optimized computations, parallel processing, and efficient memory management, which are crucial for training large-scale machine learning models.
- **Building Blocks of Computation Graphs:** In TensorFlow's graph-based execution model (prior to eager execution becoming default), tensors represented the data flowing through the nodes of a computational graph. Each operation in the graph takes tensors as input and produces tensors as output, forming a directed acyclic graph that defines the entire computation.
- **Foundation for Machine Learning Operations:** Nearly every operation in a machine learning model built with TensorFlow, from simple arithmetic to complex neural network layers (like convolutions, activations, and pooling), is defined and executed as a tensor operation.
- **Scalability and Flexibility:** The tensor concept allows TensorFlow to scale computations across different devices and platforms. Whether running on a CPU, GPU, or TPU, the underlying data structure remains a tensor, ensuring consistency and enabling the framework to leverage hardware acceleration effectively.

## Q7. How can TensorBoard be integrated with Tensorflow 2.0?

Ans-> TensorBoard can be integrated with TensorFlow 2.0 primarily through the use of Keras Callbacks and tf.summary API for custom logging.

1. Using the Keras TensorBoard Callback: This is the most common and straightforward method for integrating TensorBoard with Keras models in TensorFlow 2.0. Import the Callback. Python

from tensorflow.keras.callbacks import TensorBoard Instantiate the Callback. Python

log\_dir = "logs/fit/" # Define a directory to store TensorBoard logs tensorboard\_callback = TensorBoard(log\_dir=log\_dir, histogram\_freq=1) # histogram\_freq logs histograms every 'n' epochs Add to Model Training. Pass the tensorboard\_callback to the callbacks argument of your model.fit() method: Python

model.fit(x\_train, y\_train, epochs=10, callbacks=[tensorboard\_callback]) This automatically logs metrics like loss and accuracy, as well as model graphs and histograms of weights/biases.

2. Using the tf.summary API for Custom Logging: For logging custom metrics or data not directly handled by the Keras callback, you can use the tf.summary API. Create a Summary File Writer. Python

logdir = "logs/custom\_logs/" writer = tf.summary.create\_file\_writer(logdir) Log Custom Data within a with Block: Use tf.summary.scalar(), tf.summary.image(), tf.summary.histogram(), etc., within a with writer.as\_default(): block: Python

```
tf.summary.scalar('custom_metric', my_custom_value, step=epoch_number)
```

3. Launching TensorBoard: After generating logs, launch TensorBoard from your terminal or within a Jupyter/Colab notebook. From the Command Line. Code

tensorboard --logdir logs/ (Replace logs/ with your actual log directory.) Within a Jupyter/Colab Notebook. Python

%load\_ext tensorboard %tensorboard --logdir logs/ TensorBoard will then launch in your web browser (usually at <http://localhost:6006>), allowing you to visualize your training progress, model graphs, and other logged data.

## Q8. What is the purpose of Tensorflow playground?

Ans-> TensorFlow Playground is a web-based, interactive tool designed to help users visualize and understand how neural networks function. Its primary purpose is to provide a hands-on learning environment for exploring fundamental concepts of machine learning and deep learning without requiring any coding.

- Key purposes of TensorFlow Playground include:
- **Visualizing Neural Network Behavior:** It allows users to observe in real-time how changes to network architecture (number of layers, neurons), activation functions, regularization, and other hyperparameters affect the model's performance and decision boundaries.
- **Building Intuition about Neural Networks:** By directly manipulating parameters and seeing immediate results, users can develop an intuitive understanding of concepts like learning rate, overfitting, underfitting, and the role of different features.
- **Experimentation and Exploration:** Users can experiment with various datasets, add noise, apply different feature transformations, and adjust training parameters to see how these factors influence the learning process and the final model.
- **Accessibility to Deep Learning:** It makes complex deep learning concepts more accessible to beginners by providing a visual and interactive interface, removing the need for programming knowledge or complex mathematical formulas.

## Q9. What is Netron, and how is it useful for deep learning models?

Ans-> This deep neural network is trained and validated on simulated database, and tested on real scattering images. The fundamental properties of the neutron make it a powerful tool to investigate atomic-scale structure and dynamics of materials.

## Q10. What is the difference between TensorFlow and Pytorch?

Ans-> TensorFlow and PyTorch are both open-source deep learning frameworks widely used for building and training neural networks, but they differ in several key aspects:

### 1. Computation Graphs:

- PyTorch: Uses dynamic computation graphs (also known as "define-by-run"). This means the graph is built on the fly as operations are executed, offering greater flexibility for debugging and research, as well as easier integration with standard Python control flow.
- TensorFlow: Traditionally used static computation graphs (also known as "define-and-run"), where the entire graph is defined before execution. While offering potential for optimization and deployment, it can be less intuitive for debugging and requires a separate compilation step. However, TensorFlow has introduced Eager Execution to provide a more PyTorch-like dynamic experience.

### 2. Ease of Use and Pythonic Nature:

- PyTorch: Is generally considered more "Pythonic" and easier to learn for those familiar with Python, as it integrates seamlessly with Python's native features and debugging tools.
- TensorFlow: Can have a steeper learning curve, especially with its static graph paradigm, although its Eager Execution mode addresses some of these concerns.

### 3. Deployment and Production Readiness:

- TensorFlow: Has a more mature and comprehensive ecosystem for production deployment, including tools like TensorFlow Serving for efficient model serving and TensorFlow Lite for mobile and edge devices.
- PyTorch: While improving significantly, its deployment ecosystem is still evolving compared to TensorFlow's established solutions.

### 4. Community and Adoption:

- TensorFlow: Being older and backed by Google, it traditionally had a larger community and broader industry adoption, particularly in large-scale production environments.
- PyTorch: Has seen rapid growth in popularity, especially within the research community, due to its flexibility and ease of use for experimentation.

### 5. Debugging and Flexibility:

- PyTorch: Its dynamic graphs make debugging more straightforward as operations are executed immediately, allowing for direct inspection of intermediate values.
- TensorFlow: Debugging static graphs can be more challenging due to the compilation step, though Eager Execution helps bridge this gap.
- In essence, PyTorch is often favored for research and rapid prototyping due to its flexibility and Pythonic nature, while TensorFlow remains a strong choice for large-scale production deployments and established industrial applications, though both frameworks are constantly evolving and incorporating features from each other.

## Q11. How do you install PyTorch?

Ans-> Installing PyTorch typically involves using a package manager like pip or conda within a Python environment. The specific command depends on your operating system, preferred package manager, and whether you require GPU support (CUDA).

- Here's a general outline of the installation process: Set up a Python Environment: It is highly recommended to use a virtual environment (like venv or conda environments) to manage your project dependencies and avoid conflicts. Using venv: Code

```
python -m venv my_pytorch_env
source my_pytorch_env/bin/activate # On Windows: my_pytorch_env\Scripts\activate
```

Using Anaconda/Miniconda. Code

```
conda create -n my_pytorch_env python=3.x
conda activate my_pytorch_env
```

## Q12. What is the basic structure of a PyTorch neural network?

Ans-> A PyTorch neural network is fundamentally built by subclassing `torch.nn.Module`. This class provides the core functionality for defining and managing neural network layers and their parameters.

- The basic structure involves: Defining the Network Class: Create a class that inherits from `nn.Module`. Initialize the network's layers within the `__init__` method of this class. These layers are typically instances of `nn.Linear` (for fully connected layers), `nn.Conv2d` (for

convolutional layers), nn.ReLU (for activation functions), etc. Python

```
import torch.nn as nn
```

```
class MyNeuralNetwork(nn.Module):
```

```
    def __init__(self):
        super().__init__()
        self.fc1 = nn.Linear(in_features=10, out_features=5)
        self.relu = nn.ReLU()
        self.fc2 = nn.Linear(in_features=5, out_features=1)
```

- Implementing the Forward Pass: Define the forward method within the network class. This method specifies how input data flows through the defined layers and operations to produce an output. This is where the computational graph of the network is implicitly built. Python

```
class MyNeuralNetwork(nn.Module):
```

```
    def __init__(self):
        super().__init__()
        self.fc1 = nn.Linear(in_features=10, out_features=5)
        self.relu = nn.ReLU()
        self.fc2 = nn.Linear(in_features=5, out_features=1)
```

```
    def forward(self, x):
        x = self.fc1(x)
        x = self.relu(x)
        x = self.fc2(x)
        return x
```

- Instantiation and Usage: Create an instance of your network class. Pass input tensors to the network instance to perform a forward pass and obtain predictions. Python

```
import torch
```

```
model = MyNeuralNetwork() input_tensor = torch.randn(1, 10) # Example input with 10 features output = model(input_tensor) This structure allows for modularity, easy management of parameters, and automatic differentiation during training, which are key features of PyTorch.
```

### Q13. What is the significance of tensors in PyTorch?

Ans-> Tensors are the fundamental data structure in PyTorch, serving as the cornerstone for all operations within the framework. Their significance stems from several key aspects:

- Data Representation: Tensors are multi-dimensional arrays, analogous to NumPy's ndarrays, used to encode the inputs, outputs, and parameters of a machine learning model. This includes representing various data types like images (as 3D or 4D tensors), text (as sequences of word embeddings), audio, and more.
- GPU Acceleration: A critical advantage of PyTorch tensors over NumPy arrays is their ability to leverage GPUs for accelerated computation. Tensors can be seamlessly moved between CPU and GPU memory, enabling significantly faster training and inference for deep learning models, especially with large datasets and complex architectures.
- Automatic Differentiation (Autograd): Tensors are integrated with PyTorch's autograd engine, which automatically computes gradients for all operations performed on them. This is crucial for backpropagation in neural networks, allowing for efficient optimization of model parameters during training.
- Foundation for Operations: All mathematical and logical operations within PyTorch, from basic arithmetic to complex neural network layers, are designed to operate on tensors. This provides a consistent and efficient interface for building and manipulating computational graphs.
- Flexibility and Interoperability: PyTorch tensors offer a high degree of flexibility in terms of data types, shapes, and device allocation. They also provide a bridge to NumPy, allowing for easy conversion between PyTorch tensors and NumPy arrays when needed.

### Q14. What is the difference between torch.Tensor and torch.cuda.Tensor in Pytorch?

Ans-> The core difference between torch.Tensor and torch.cuda.Tensor in PyTorch lies in the device where their data is stored and where computations are performed.

- torch.Tensor: This refers to a tensor whose data resides in CPU memory. Operations on torch.Tensor objects are executed by the CPU. This is the default tensor type when you create a tensor without explicitly specifying a device.
- torch.cuda.Tensor: This refers to a tensor whose data resides in GPU memory. Operations on torch.cuda.Tensor objects are executed by the GPU. You typically obtain a torch.cuda.Tensor by moving a torch.Tensor to a CUDA-enabled GPU using methods like .to('cuda') or .cuda().

- Key differences summarized:
- Memory Location: `torch.Tensor` is on the CPU, `torch.cuda.Tensor` is on the GPU.
- Computational Device: Operations on `torch.Tensor` use the CPU, while operations on `torch.cuda.Tensor` use the GPU.
- Performance: For computationally intensive tasks, `torch.cuda.Tensor` generally offers significantly faster execution due to the parallel processing capabilities of GPUs.
- Interoperability: You cannot directly perform operations that mix `torch.Tensor` and `torch.cuda.Tensor` objects. All tensors involved in an operation must be on the same device. You must explicitly move tensors between CPU and GPU if needed.

### Q15. What is the purpose of the `torch.optim` module in PyTorch?

Ans-> The `torch.optim` module in PyTorch serves the purpose of providing a collection of optimization algorithms commonly used for training neural networks. These algorithms are essential for minimizing the loss function of a model by iteratively adjusting its parameters (weights and biases) based on the computed gradients.

- Here's a breakdown of its key functions: Implementing Optimization Algorithms: `torch.optim` offers various optimization algorithms, including popular ones like Stochastic Gradient Descent (SGD), Adam, RMSprop, and more. Each algorithm has its own strategy for updating parameters to navigate the loss landscape effectively.
- Managing Model Parameters: Optimizers within `torch.optim` are initialized by providing them with the parameters of the model to be optimized. These parameters are typically obtained using `model.parameters()`, which returns an iterable of `torch.nn.Parameter` objects.
- Updating Parameters Based on Gradients: After a forward pass and backward pass (which computes gradients), the optimizer's `step()` method is called. This method uses the calculated gradients and the chosen optimization algorithm to update the model's parameters, moving them in a direction that reduces the loss.
- Handling Optimizer State: The optimizer objects maintain internal state (e.g., momentum buffers in Adam) that is crucial for their operation. This state is managed automatically by the `torch.optim` module.
- Zeroing Gradients: Before computing gradients for a new batch of data, the optimizer's `zero_grad()` method is called to reset the gradients of all optimized parameters to zero. This prevents gradients from accumulating across iterations.
- Supporting Advanced Features: `torch.optim` also supports advanced features like learning rate scheduling (adjusting the learning rate during training) and weight decay (a regularization technique to prevent overfitting).

### Q16. What are some common activation functions used in neural networks?

Ans-> Common activation functions in neural networks include ReLU, sigmoid, tanh, and softmax. These functions introduce non-linearity, enabling the network to learn complex patterns in data. ReLU is widely used, particularly in convolutional neural networks, while sigmoid and tanh are often used in earlier layers. Softmax is typically used in the output layer for multi-class classification.

- Here's a more detailed look: ReLU (Rectified Linear Unit):  $f(x) = \max(0, x)$ . It outputs the input directly if it's positive, and zero otherwise. ReLU is computationally efficient and helps mitigate the vanishing gradient problem during training.
- Sigmoid:  $f(x) = 1 / (1 + \exp(-x))$ . It outputs a value between 0 and 1. Sigmoid was historically popular but is prone to the vanishing gradient problem, especially in deep networks.
- Tanh (Hyperbolic Tangent):  $f(x) = (\exp(x) - \exp(-x)) / (\exp(x) + \exp(-x))$ . It outputs a value between -1 and 1. Tanh is also susceptible to the vanishing gradient problem, but generally performs better than sigmoid in some cases.
- Softmax: Used in the output layer for multi-class classification. It normalizes the output of each neuron into a probability distribution, where the sum of all probabilities equals 1.

### Q17. What is the difference between `torch.nn.Module` and `torch.nn.Sequential` in Pytorch?

Ans-> `torch.nn.Module` is the base class for all neural network modules in PyTorch. Any custom layer, model, or component that you create in PyTorch should inherit from `torch.nn.Module`. This base class provides essential functionalities like tracking parameters, managing submodules, and providing the `forward()` method where the computation logic is defined. `torch.nn.Sequential` is a specific type of `torch.nn.Module` that acts as a container for other modules. It allows you to stack multiple modules in a sequential order, where the output of one module becomes the input of the next. This simplifies the creation of feed-forward networks by automatically handling the forward pass through the sequence of modules.

- Key Differences: Generality vs. Specificity: `torch.nn.Module` is a general base class for any neural network component, while `torch.nn.Sequential` is a specialized container for modules that are meant to be executed in a specific, linear order.
- Flexibility: When inheriting from `torch.nn.Module`, you have full control over the `forward()` method, allowing for complex and non-sequential network architectures (e.g., skip connections, multiple inputs/outputs). `torch.nn.Sequential` is designed for simple, sequential data flow.
- Usage: You define custom layers or entire models by subclassing `torch.nn.Module` and implementing its `forward()` method. You use `torch.nn.Sequential` to easily chain together existing modules or custom `nn.Module` instances into a single, sequential unit.

### Q18. How can you monitor training progress in Tensorflow 2.0?

Ans-> Monitoring training progress in TensorFlow 2.0 can be effectively achieved using several methods:

1. TensorBoard: TensorBoard is TensorFlow's visualization toolkit, providing a web-based interface for monitoring various aspects of training:
  - Metrics: Visualize loss, accuracy, and other custom metrics over time.
  - Graphs: Inspect the computational graph of your model.
  - Histograms: View the distribution of weights, biases, and activations across different layers.
  - Images/Audio/Text: Visualize samples from your dataset or model outputs.
  - Profiler: Analyze model performance and identify bottlenecks. To use TensorBoard, add the `tf.keras.callbacks.TensorBoard` callback to your `model.fit()` call and specify a `log_dir`. Then, launch TensorBoard from your terminal, pointing it to the same `log_dir`.
2. Keras Callbacks: TensorFlow 2.0 leverages Keras for high-level API usage, and Keras offers various built-in callbacks to monitor and control training:
  - ModelCheckpoint: Saves the model at regular intervals or when a specific metric improves.
  - EarlyStopping: Stops training when a monitored metric stops improving for a specified number of epochs. `ReduceLROnPlateau`: Reduces the learning rate when a metric has stopped improving.
  - CSVLogger: Streams epoch results to a CSV file.
3. Custom Training Loops with `tf.keras.utils.Progbar`: For custom training loops, `tf.keras.utils.Progbar` can be used to display a progress bar in the console, providing visual feedback on the iteration progress within an epoch.
4. Manual Printing: For basic monitoring, you can simply print out metrics like loss and accuracy at the end of each epoch within your training loop. This offers immediate feedback in the console.
5. TensorFlow Profiler: For in-depth performance analysis and bottleneck identification, the TensorFlow Profiler, integrated within TensorBoard, provides detailed insights into resource utilization (CPU, GPU) and operation execution times during training.

## Q19. How does the Keras API fit into Tensorflow 2.0?

Ans-> In TensorFlow 2.0, Keras is integrated as the official high-level API for building and training deep learning models. This means that Keras is no longer a separate library that uses TensorFlow as a backend; instead, it is a core component of TensorFlow itself, accessible through the `tf.keras` module.

- This integration provides several benefits:
- Simplified API: Keras offers a user-friendly and intuitive API that abstracts away the complexities of low-level TensorFlow operations, making it easier to build and experiment with neural networks.
- Seamless Integration: `tf.keras` is tightly integrated with other TensorFlow functionalities, including `tf.data` for efficient data pipelines, `tf.distribute` for distributed training, and `tf.saved_model` for model deployment.
- Eager Execution Support: Keras in TensorFlow 2.0 fully supports eager execution, allowing for more interactive and flexible development and debugging.
- Standardization: By embracing Keras, TensorFlow 2.0 provides a standardized and recommended way to build deep learning models, promoting consistency and ease of collaboration within the TensorFlow ecosystem.
- While the standalone Keras library (accessible via `import keras`) still exists and can be used with various backends, `tf.keras` is the preferred and fully optimized version for use within TensorFlow 2.0 and beyond. This integration streamlines the development process and leverages the full power of the TensorFlow framework.

## Q20. What is the example of deep learning project that can be implemented using Tensorflow 2.0?

Ans-> A common example of a deep learning project implementable with TensorFlow 2.0 is image classification using Convolutional Neural Networks (CNNs).

- This project involves training a model to categorize images into predefined classes. For instance, one could build a system to classify images of different animals, objects, or even handwritten digits (like the MNIST dataset).
- Key steps in implementing this project with TensorFlow 2.0:
  - Data Preparation: Load and preprocess a dataset of images, ensuring they are correctly labeled and resized to a consistent input shape.
  - Model Architecture: Define a CNN architecture using TensorFlow's Keras API. This typically involves layers like `Conv2D`, `MaxPooling2D`, `Flatten`, and `Dense`.
  - Compilation: Compile the model by specifying an optimizer (e.g., Adam), a loss function (e.g., `categorical_crossentropy` for multi-class classification), and metrics (e.g., accuracy).
  - Training: Train the model on the prepared dataset using the `model.fit()` method.
  - Evaluation: Evaluate the trained model's performance on a separate test set to assess its accuracy and generalization capabilities.
- Prediction: Use the trained model to make predictions on new, unseen images. TensorFlow 2.0's eager execution and simplified API make it convenient to build and experiment with such deep learning models.

## Q21. What is the main advantage of using pre-trained models in Tensorflow and PyTorch?

Ans-> The main advantage of using pre-trained models in TensorFlow and PyTorch is transfer learning, which significantly reduces the need for extensive data and computational resources while accelerating development.

- Pre-trained models have already learned general features from vast datasets, such as ImageNet for computer vision or large text corpora for natural language processing. Instead of training a model from scratch, which requires massive amounts of data and

computational power, these pre-trained models can be fine-tuned on smaller, domain-specific datasets for new tasks. This allows for:

- **Reduced Training Time and Computational Cost:** Leveraging pre-existing learned features eliminates the need for lengthy initial training phases.
  - **Improved Performance with Limited Data:** Pre-trained models can achieve high accuracy even with relatively small custom datasets, as they already possess a strong foundation of knowledge.
- Access to State-of-the-Art Architectures:** Utilizing pre-trained models provides access to complex and well-validated architectures without the need for extensive design and testing.

## Practical

### Q1. How do you install and verify that Tensorflow 2.0 was installed successfully?

```
# Install a compatible version of TensorFlow
!pip install tensorflow==2.10
```

```
# Verify the installation
import tensorflow as tf
print(tf.__version__)
```

```
ERROR: Could not find a version that satisfies the requirement tensorflow==2.10 (from versions: 2.12.0rc0, 2.12.0rc1, 2.12.0, 2.12.1, 2.12.2, 2.12.3, 2.12.4, 2.12.5, 2.12.6, 2.12.7, 2.12.8, 2.12.9, 2.12.10, 2.12.11, 2.12.12, 2.12.13, 2.12.14, 2.12.15, 2.12.16, 2.12.17, 2.12.18, 2.12.19, 2.12.20, 2.12.21, 2.12.22, 2.12.23, 2.12.24, 2.12.25, 2.12.26, 2.12.27, 2.12.28, 2.12.29, 2.12.30, 2.12.31, 2.12.32, 2.12.33, 2.12.34, 2.12.35, 2.12.36, 2.12.37, 2.12.38, 2.12.39, 2.12.40, 2.12.41, 2.12.42, 2.12.43, 2.12.44, 2.12.45, 2.12.46, 2.12.47, 2.12.48, 2.12.49, 2.12.50, 2.12.51, 2.12.52, 2.12.53, 2.12.54, 2.12.55, 2.12.56, 2.12.57, 2.12.58, 2.12.59, 2.12.60, 2.12.61, 2.12.62, 2.12.63, 2.12.64, 2.12.65, 2.12.66, 2.12.67, 2.12.68, 2.12.69, 2.12.70, 2.12.71, 2.12.72, 2.12.73, 2.12.74, 2.12.75, 2.12.76, 2.12.77, 2.12.78, 2.12.79, 2.12.80, 2.12.81, 2.12.82, 2.12.83, 2.12.84, 2.12.85, 2.12.86, 2.12.87, 2.12.88, 2.12.89, 2.12.90, 2.12.91, 2.12.92, 2.12.93, 2.12.94, 2.12.95, 2.12.96, 2.12.97, 2.12.98, 2.12.99, 2.12.100, 2.12.101, 2.12.102, 2.12.103, 2.12.104, 2.12.105, 2.12.106, 2.12.107, 2.12.108, 2.12.109, 2.12.110, 2.12.111, 2.12.112, 2.12.113, 2.12.114, 2.12.115, 2.12.116, 2.12.117, 2.12.118, 2.12.119, 2.12.120, 2.12.121, 2.12.122, 2.12.123, 2.12.124, 2.12.125, 2.12.126, 2.12.127, 2.12.128, 2.12.129, 2.12.130, 2.12.131, 2.12.132, 2.12.133, 2.12.134, 2.12.135, 2.12.136, 2.12.137, 2.12.138, 2.12.139, 2.12.140, 2.12.141, 2.12.142, 2.12.143, 2.12.144, 2.12.145, 2.12.146, 2.12.147, 2.12.148, 2.12.149, 2.12.150, 2.12.151, 2.12.152, 2.12.153, 2.12.154, 2.12.155, 2.12.156, 2.12.157, 2.12.158, 2.12.159, 2.12.160, 2.12.161, 2.12.162, 2.12.163, 2.12.164, 2.12.165, 2.12.166, 2.12.167, 2.12.168, 2.12.169, 2.12.170, 2.12.171, 2.12.172, 2.12.173, 2.12.174, 2.12.175, 2.12.176, 2.12.177, 2.12.178, 2.12.179, 2.12.180, 2.12.181, 2.12.182, 2.12.183, 2.12.184, 2.12.185, 2.12.186, 2.12.187, 2.12.188, 2.12.189, 2.12.190, 2.12.191, 2.12.192, 2.12.193, 2.12.194, 2.12.195, 2.12.196, 2.12.197, 2.12.198, 2.12.199, 2.12.200, 2.12.201, 2.12.202, 2.12.203, 2.12.204, 2.12.205, 2.12.206, 2.12.207, 2.12.208, 2.12.209, 2.12.210, 2.12.211, 2.12.212, 2.12.213, 2.12.214, 2.12.215, 2.12.216, 2.12.217, 2.12.218, 2.12.219, 2.12.220, 2.12.221, 2.12.222, 2.12.223, 2.12.224, 2.12.225, 2.12.226, 2.12.227, 2.12.228, 2.12.229, 2.12.230, 2.12.231, 2.12.232, 2.12.233, 2.12.234, 2.12.235, 2.12.236, 2.12.237, 2.12.238, 2.12.239, 2.12.240, 2.12.241, 2.12.242, 2.12.243, 2.12.244, 2.12.245, 2.12.246, 2.12.247, 2.12.248, 2.12.249, 2.12.250, 2.12.251, 2.12.252, 2.12.253, 2.12.254, 2.12.255, 2.12.256, 2.12.257, 2.12.258, 2.12.259, 2.12.260, 2.12.261, 2.12.262, 2.12.263, 2.12.264, 2.12.265, 2.12.266, 2.12.267, 2.12.268, 2.12.269, 2.12.270, 2.12.271, 2.12.272, 2.12.273, 2.12.274, 2.12.275, 2.12.276, 2.12.277, 2.12.278, 2.12.279, 2.12.280, 2.12.281, 2.12.282, 2.12.283, 2.12.284, 2.12.285, 2.12.286, 2.12.287, 2.12.288, 2.12.289, 2.12.290, 2.12.291, 2.12.292, 2.12.293, 2.12.294, 2.12.295, 2.12.296, 2.12.297, 2.12.298, 2.12.299, 2.12.300, 2.12.301, 2.12.302, 2.12.303, 2.12.304, 2.12.305, 2.12.306, 2.12.307, 2.12.308, 2.12.309, 2.12.310, 2.12.311, 2.12.312, 2.12.313, 2.12.314, 2.12.315, 2.12.316, 2.12.317, 2.12.318, 2.12.319, 2.12.320, 2.12.321, 2.12.322, 2.12.323, 2.12.324, 2.12.325, 2.12.326, 2.12.327, 2.12.328, 2.12.329, 2.12.330, 2.12.331, 2.12.332, 2.12.333, 2.12.334, 2.12.335, 2.12.336, 2.12.337, 2.12.338, 2.12.339, 2.12.340, 2.12.341, 2.12.342, 2.12.343, 2.12.344, 2.12.345, 2.12.346, 2.12.347, 2.12.348, 2.12.349, 2.12.350, 2.12.351, 2.12.352, 2.12.353, 2.12.354, 2.12.355, 2.12.356, 2.12.357, 2.12.358, 2.12.359, 2.12.360, 2.12.361, 2.12.362, 2.12.363, 2.12.364, 2.12.365, 2.12.366, 2.12.367, 2.12.368, 2.12.369, 2.12.370, 2.12.371, 2.12.372, 2.12.373, 2.12.374, 2.12.375, 2.12.376, 2.12.377, 2.12.378, 2.12.379, 2.12.380, 2.12.381, 2.12.382, 2.12.383, 2.12.384, 2.12.385, 2.12.386, 2.12.387, 2.12.388, 2.12.389, 2.12.390, 2.12.391, 2.12.392, 2.12.393, 2.12.394, 2.12.395, 2.12.396, 2.12.397, 2.12.398, 2.12.399, 2.12.400, 2.12.401, 2.12.402, 2.12.403, 2.12.404, 2.12.405, 2.12.406, 2.12.407, 2.12.408, 2.12.409, 2.12.410, 2.12.411, 2.12.412, 2.12.413, 2.12.414, 2.12.415, 2.12.416, 2.12.417, 2.12.418, 2.12.419, 2.12.420, 2.12.421, 2.12.422, 2.12.423, 2.12.424, 2.12.425, 2.12.426, 2.12.427, 2.12.428, 2.12.429, 2.12.430, 2.12.431, 2.12.432, 2.12.433, 2.12.434, 2.12.435, 2.12.436, 2.12.437, 2.12.438, 2.12.439, 2.12.440, 2.12.441, 2.12.442, 2.12.443, 2.12.444, 2.12.445, 2.12.446, 2.12.447, 2.12.448, 2.12.449, 2.12.450, 2.12.451, 2.12.452, 2.12.453, 2.12.454, 2.12.455, 2.12.456, 2.12.457, 2.12.458, 2.12.459, 2.12.460, 2.12.461, 2.12.462, 2.12.463, 2.12.464, 2.12.465, 2.12.466, 2.12.467, 2.12.468, 2.12.469, 2.12.470, 2.12.471, 2.12.472, 2.12.473, 2.12.474, 2.12.475, 2.12.476, 2.12.477, 2.12.478, 2.12.479, 2.12.480, 2.12.481, 2.12.482, 2.12.483, 2.12.484, 2.12.485, 2.12.486, 2.12.487, 2.12.488, 2.12.489, 2.12.490, 2.12.491, 2.12.492, 2.12.493, 2.12.494, 2.12.495, 2.12.496, 2.12.497, 2.12.498, 2.12.499, 2.12.500, 2.12.501, 2.12.502, 2.12.503, 2.12.504, 2.12.505, 2.12.506, 2.12.507, 2.12.508, 2.12.509, 2.12.510, 2.12.511, 2.12.512, 2.12.513, 2.12.514, 2.12.515, 2.12.516, 2.12.517, 2.12.518, 2.12.519, 2.12.520, 2.12.521, 2.12.522, 2.12.523, 2.12.524, 2.12.525, 2.12.526, 2.12.527, 2.12.528, 2.12.529, 2.12.530, 2.12.531, 2.12.532, 2.12.533, 2.12.534, 2.12.535, 2.12.536, 2.12.537, 2.12.538, 2.12.539, 2.12.540, 2.12.541, 2.12.542, 2.12.543, 2.12.544, 2.12.545, 2.12.546, 2.12.547, 2.12.548, 2.12.549, 2.12.550, 2.12.551, 2.12.552, 2.12.553, 2.12.554, 2.12.555, 2.12.556, 2.12.557, 2.12.558, 2.12.559, 2.12.560, 2.12.561, 2.12.562, 2.12.563, 2.12.564, 2.12.565, 2.12.566, 2.12.567, 2.12.568, 2.12.569, 2.12.570, 2.12.571, 2.12.572, 2.12.573, 2.12.574, 2.12.575, 2.12.576, 2.12.577, 2.12.578, 2.12.579, 2.12.580, 2.12.581, 2.12.582, 2.12.583, 2.12.584, 2.12.585, 2.12.586, 2.12.587, 2.12.588, 2.12.589, 2.12.590, 2.12.591, 2.12.592, 2.12.593, 2.12.594, 2.12.595, 2.12.596, 2.12.597, 2.12.598, 2.12.599, 2.12.600, 2.12.601, 2.12.602, 2.12.603, 2.12.604, 2.12.605, 2.12.606, 2.12.607, 2.12.608, 2.12.609, 2.12.610, 2.12.611, 2.12.612, 2.12.613, 2.12.614, 2.12.615, 2.12.616, 2.12.617, 2.12.618, 2.12.619, 2.12.620, 2.12.621, 2.12.622, 2.12.623, 2.12.624, 2.12.625, 2.12.626, 2.12.627, 2.12.628, 2.12.629, 2.12.630, 2.12.631, 2.12.632, 2.12.633, 2.12.634, 2.12.635, 2.12.636, 2.12.637, 2.12.638, 2.12.639, 2.12.640, 2.12.641, 2.12.642, 2.12.643, 2.12.644, 2.12.645, 2.12.646, 2.12.647, 2.12.648, 2.12.649, 2.12.650, 2.12.651, 2.12.652, 2.12.653, 2.12.654, 2.12.655, 2.12.656, 2.12.657, 2.12.658, 2.12.659, 2.12.660, 2.12.661, 2.12.662, 2.12.663, 2.12.664, 2.12.665, 2.12.666, 2.12.667, 2.12.668, 2.12.669, 2.12.670, 2.12.671, 2.12.672, 2.12.673, 2.12.674, 2.12.675, 2.12.676, 2.12.677, 2.12.678, 2.12.679, 2.12.680, 2.12.681, 2.12.682, 2.12.683, 2.12.684, 2.12.685, 2.12.686, 2.12.687, 2.12.688, 2.12.689, 2.12.690, 2.12.691, 2.12.692, 2.12.693, 2.12.694, 2.12.695, 2.12.696, 2.12.697, 2.12.698, 2.12.699, 2.12.700, 2.12.701, 2.12.702, 2.12.703, 2.12.704, 2.12.705, 2.12.706, 2.12.707, 2.12.708, 2.12.709, 2.12.710, 2.12.711, 2.12.712, 2.12.713, 2.12.714, 2.12.715, 2.12.716, 2.12.717, 2.12.718, 2.12.719, 2.12.720, 2.12.721, 2.12.722, 2.12.723, 2.12.724, 2.12.725, 2.12.726, 2.12.727, 2.12.728, 2.12.729, 2.12.730, 2.12.731, 2.12.732, 2.12.733, 2.12.734, 2.12.735, 2.12.736, 2.12.737, 2.12.738, 2.12.739, 2.12.740, 2.12.741, 2.12.742, 2.12.743, 2.12.744, 2.12.745, 2.12.746, 2.12.747, 2.12.748, 2.12.749, 2.12.750, 2.12.751, 2.12.752, 2.12.753, 2.12.754, 2.12.755, 2.12.756, 2.12.757, 2.12.758, 2.12.759, 2.12.760, 2.12.761, 2.12.762, 2.12.763, 2.12.764, 2.12.765, 2.12.766, 2.12.767, 2.12.768, 2.12.769, 2.12.770, 2.12.771, 2.12.772, 2.12.773, 2.12.774, 2.12.775, 2.12.776, 2.12.777, 2.12.778, 2.12.779, 2.12.780, 2.12.781, 2.12.782, 2.12.783, 2.12.784, 2.12.785, 2.12.786, 2.12.787, 2.12.788, 2.12.789, 2.12.790, 2.12.791, 2.12.792, 2.12.793, 2.12.794, 2.12.795, 2.12.796, 2.12.797, 2.12.798, 2.12.799, 2.12.800, 2.12.801, 2.12.802, 2.12.803, 2.12.804, 2.12.805, 2.12.806, 2.12.807, 2.12.808, 2.12.809, 2.12.810, 2.12.811, 2.12.812, 2.12.813, 2.12.814, 2.12.815, 2.12.816, 2.12.817, 2.12.818, 2.12.819, 2.12.820, 2.12.821, 2.12.822, 2.12.823, 2.12.824, 2.12.825, 2.12.826, 2.12.827, 2.12.828, 2.12.829, 2.12.830, 2.12.831, 2.12.832, 2.12.833, 2.12.834, 2.12.835, 2.12.836, 2.12.837, 2.12.838, 2.12.839, 2.12.840, 2.12.841, 2.12.842, 2.12.843, 2.12.844, 2.12.845, 2.12.846, 2.12.847, 2.12.848, 2.12.849, 2.12.850, 2.12.851, 2.12.852, 2.12.853, 2.12.854, 2.12.855, 2.12.856, 2.12.857, 2.12.858, 2.12.859, 2.12.860, 2.12.861, 2.12.862, 2.12.863, 2.12.864, 2.12.865, 2.12.866, 2.12.867, 2.12.868, 2.12.869, 2.12.870, 2.12.871, 2.12.872, 2.12.873, 2.12.874, 2.12.875, 2.12.876, 2.12.877, 2.12.878, 2.12.879, 2.12.880, 2.12.881, 2.12.882, 2.12.883, 2.12.884, 2.12.885, 2.12.886, 2.12.887, 2.12.888, 2.12.889, 2.12.890, 2.12.891, 2.12.892, 2.12.893, 2.12.894, 2.12.895, 2.12.896, 2.12.897, 2.12.898, 2.12.899, 2.12.900, 2.12.901, 2.12.902, 2.12.903, 2.12.904, 2.12.905, 2.12.906, 2.12.907, 2.12.908, 2.12.909, 2.12.910, 2.12.911, 2.12.912, 2.12.913, 2.12.914, 2.12.915, 2.12.916, 2.12.917, 2.12.918, 2.12.919, 2.12.920, 2.12.921, 2.12.922, 2.12.923, 2.12.924, 2.12.925, 2.12.926, 2.12.927, 2.12.928, 2.12.929, 2.12.930, 2.12.931, 2.12.932, 2.12.933, 2.12.934, 2.12.935, 2.12.936, 2.12.937, 2.12.938, 2.12.939, 2.12.940, 2.12.941, 2.12.942, 2.12.943, 2.12.944, 2.12.945, 2.12.946, 2.12.947, 2.12.948, 2.12.949, 2.12.950, 2.12.951, 2.12.952, 2.12.953, 2.12.954, 2.12.955, 2.12.956, 2.12.957, 2.12.958, 2.12.959, 2.12.960, 2.12.961, 2.12.962, 2.12.963, 2.12.964, 2.12.965, 2.12.966, 2.12.967, 2.12.968, 2.12.969, 2.12.970, 2.12.971, 2.12.972, 2.12.973, 2.12.974, 2.12.975, 2.12.976, 2.12.977, 2.12.978, 2.12.979, 2.12.980, 2.12.981, 2.12.982, 2.12.983, 2.12.984, 2.12.985, 2.12.986, 2.12.987, 2.12.988, 2.12.989, 2.12.990, 2.12.991, 2.12.992, 2.12.993, 2.12.994, 2.12.995, 2.12.996, 2.12.997, 2.12.998, 2.12.999, 2.13.0, 2.13.1, 2.13.2, 2.13.3, 2.13.4, 2.13.5, 2.13.6, 2.13.7, 2.13.8, 2.13.9, 2.13.10, 2.13.11, 2.13.12, 2.13.13, 2.13.14, 2.13.15, 2.13.16, 2.13.17, 2.13.18, 2.13.19, 2.13.20, 2.13.21, 2.13.22, 2.13.23, 2.13.24, 2.13.25, 2.13.26, 2.13.27, 2.13.28, 2.13.29, 2.13.30, 2.13.31, 2.13.32, 2.13.33, 2.13.34, 2.13.35, 2.13.36, 2.13.37, 2.13.38, 2.13.39, 2.13.40, 2.13.41, 2.13.42, 2.13.43, 2.13.44, 2.13.45, 2.13.46, 2.13.47, 2.13.48, 2.13.49, 2.13.50, 2.13.51, 2.13.52, 2.13.53, 2.13.54, 2.13.55, 2.13.56, 2.13.57, 2.13.58, 2.13.59, 2.13.60, 2.13.61, 2.13.62, 2.13.63, 2.13.64, 2.13.65, 2.13.66, 2.13.67, 2.13.68, 2.13.69, 2.13.70, 2.13.71, 2.13.72, 2.13.73, 2.13.74, 2.13.75, 2.13.76, 2.13.77, 2.13.78, 2.13.79, 2.13.80, 2.13.81, 2.13.82, 2.13.83, 2.13.84, 2.13.85, 2.13.86, 2.13.87, 2.13.88, 2.13.89, 2.13.90, 2.13.91, 2.13.92, 2.13.93, 2.13.94, 2.13.95, 2.13.96, 2.13.97, 2.13.98, 2.13.99, 2.14.0, 2.14.1, 2.14.2, 2.14.3, 2.14.4, 2.14.5, 2.14.6, 2.14.7, 2.14.8, 2.14.9, 2.14.10, 2.14.11, 2.14.12, 2.14.13, 2.14.14, 2.14.15, 2.14.16, 2.14.17, 2.14.18, 2.14.19, 2.14.20, 2.14.21, 2.14.22, 2.14.23, 2.14.24, 2.14.25, 2.14.26, 2.14.27, 2.14.28, 2.14.29, 2.14.30, 2.14.31, 2.14.32, 2.14.33, 2.14.34, 2.14.35, 2.14.36, 2.14.37, 2.14.38, 2.14.39, 2.14.40, 2.14.41, 2.14.42, 2.14.43, 2.14.44, 2.14.45, 2.14.46, 2.14.47, 2.14.48, 2.14.49, 2.14.50, 2.14.51, 2.14.52, 2.14.53, 2.14.54, 2.14.55, 2.14.56, 2.14.57, 2.14.58, 2.14.59, 2.14.60, 2.14.61, 2.14.62, 2.14.63, 2.14.64, 2.14.65, 2.14.66, 2.14.67, 2.14.68, 2.14.69, 2.14.70, 2.14.71, 2.14.72, 2.14.73, 2.14.74, 2.14.75, 2.14.76, 2.14.77, 2.14.78, 2.14.79, 2.14.80, 2.14.81, 2.14.82, 2.14.83, 2.14.84, 2.14.85, 2.14.86, 2.14.87, 2.14.88, 2.14.89, 2.14.90, 2.14.91, 2.14.92, 2.14.93, 2.14.94, 2.14.95, 2.14.96, 2.14.97, 2.14.98, 2.14.99, 2.15.0, 2.15.1, 2.15.2, 2.15.3, 2.15.4, 2.15.5, 2.15.6, 2.15.7, 2.15.8, 2.15.9, 2.15.10, 2.15.11, 2.15.12, 2.15.13, 2.15.14, 2.15.15, 2.15.
```



```

/usr/local/lib/python3.11/dist-packages/keras/src/layers/core/dense.py:93: UserWarning: Do not pass an `input_shape`/`input_dim` arg
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)
Model: "sequential"

```

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 64)	50,240
dense_1 (Dense)	(None, 32)	2,080
dense_2 (Dense)	(None, 10)	330

```

Total params: 52,650 (205.66 KB)
Trainable params: 52,650 (205.66 KB)
Non-trainable params: 0 (0.00 B)

```

#### ✓ Q4. How can visualize the training progress using Tensorflow and Matplotlib?

```

import tensorflow as tf
from tensorflow.keras import layers, models
import matplotlib.pyplot as plt
import numpy as np

# Define a simple sequential model (reusing the structure from the previous example)
# For demonstration, we'll use a smaller input shape and number of classes
model = models.Sequential([
    layers.Dense(units=16, activation='relu', input_shape=(10,)),
    layers.Dense(units=8, activation='relu'),
    layers.Dense(units=3, activation='softmax') # Example for 3 classes
])

# Compile the model
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# Generate some dummy data for demonstration
num_samples = 1000
input_dim = 10
num_classes = 3

x_train = np.random.rand(num_samples, input_dim).astype(np.float32)
y_train = np.random.randint(0, num_classes, num_samples)

# Train the model and store the training history
history = model.fit(x_train, y_train, epochs=50, verbose=0) # Set verbose=0 to avoid printing progress during training


# Get the loss and accuracy from the training history
loss = history.history['loss']
accuracy = history.history['accuracy']

# Plot the training loss
plt.figure(figsize=(12, 4))
plt.subplot(1, 2, 1)
plt.plot(loss)
plt.title('Training Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')

# Plot the training accuracy
plt.subplot(1, 2, 2)
plt.plot(accuracy)
plt.title('Training Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')

plt.show()

```



## Q5. How do you install PyTorch and verify the PyTorch installation?

```
# To install PyTorch, you can use pip.
# The specific command depends on whether you want GPU support (CUDA).

# Option 1: Install PyTorch with CPU support only
# !pip install torch torchvision torchaudio --index-url https://download.pytorch.org/whl/cpu

# Option 2: Install PyTorch with CUDA support (recommended if you have a compatible GPU)
# The exact command might vary depending on your CUDA version.
# You can find the specific command for your system on the official PyTorch website:
# https://pytorch.org/get-started/locally/
!pip install torch torchvision torchaudio

# Verify the PyTorch installation
import torch
print(torch.__version__)

# Check if CUDA is available (if you installed with GPU support)
if torch.cuda.is_available():
    print("CUDA is available!")
    print(f"CUDA version: {torch.version.cuda}")
    print(f"GPU name: {torch.cuda.get_device_name(0)}")
else:
    print("CUDA is not available.")
```

```
Requirement already satisfied: torch in /usr/local/lib/python3.11/dist-packages (2.6.0+cu124)
Requirement already satisfied: torchvision in /usr/local/lib/python3.11/dist-packages (0.21.0+cu124)
Requirement already satisfied: torchaudio in /usr/local/lib/python3.11/dist-packages (2.6.0+cu124)
Requirement already satisfied: filelock in /usr/local/lib/python3.11/dist-packages (from torch) (3.18.0)
Requirement already satisfied: typing-extensions>=4.10.0 in /usr/local/lib/python3.11/dist-packages (from torch) (4.14.1)
Requirement already satisfied: networkx in /usr/local/lib/python3.11/dist-packages (from torch) (3.5)
Requirement already satisfied: jinja2 in /usr/local/lib/python3.11/dist-packages (from torch) (3.1.6)
Requirement already satisfied: fsspec in /usr/local/lib/python3.11/dist-packages (from torch) (2025.3.0)
Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.4.127 in /usr/local/lib/python3.11/dist-packages (from torch) (12.4.127)
Requirement already satisfied: nvidia-cuda-runtime-cu12==12.4.127 in /usr/local/lib/python3.11/dist-packages (from torch) (12.4.127)
Requirement already satisfied: nvidia-cuda-cupti-cu12==12.4.127 in /usr/local/lib/python3.11/dist-packages (from torch) (12.4.127)
Requirement already satisfied: nvidia-cudnn-cu12==9.1.0.70 in /usr/local/lib/python3.11/dist-packages (from torch) (9.1.0.70)
Requirement already satisfied: nvidia-cublas-cu12==12.4.5.8 in /usr/local/lib/python3.11/dist-packages (from torch) (12.4.5.8)
Requirement already satisfied: nvidia-cufft-cu12==11.2.1.3 in /usr/local/lib/python3.11/dist-packages (from torch) (11.2.1.3)
Requirement already satisfied: nvidia-curand-cu12==10.3.5.147 in /usr/local/lib/python3.11/dist-packages (from torch) (10.3.5.147)
Requirement already satisfied: nvidia-cusolver-cu12==11.6.1.9 in /usr/local/lib/python3.11/dist-packages (from torch) (11.6.1.9)
Requirement already satisfied: nvidia-cusparse-cu12==12.3.1.170 in /usr/local/lib/python3.11/dist-packages (from torch) (12.3.1.170)
Requirement already satisfied: nvidia-cusparselt-cu12==0.6.2 in /usr/local/lib/python3.11/dist-packages (from torch) (0.6.2)
Requirement already satisfied: nvidia-nccl-cu12==2.21.5 in /usr/local/lib/python3.11/dist-packages (from torch) (2.21.5)
Requirement already satisfied: nvidia-nvtx-cu12==12.4.127 in /usr/local/lib/python3.11/dist-packages (from torch) (12.4.127)
Requirement already satisfied: nvidia-nvjitlink-cu12==12.4.127 in /usr/local/lib/python3.11/dist-packages (from torch) (12.4.127)
Requirement already satisfied: triton==3.2.0 in /usr/local/lib/python3.11/dist-packages (from torch) (3.2.0)
Requirement already satisfied: sympy==1.13.1 in /usr/local/lib/python3.11/dist-packages (from torch) (1.13.1)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.11/dist-packages (from sympy==1.13.1->torch) (1.3.0)
Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages (from torchvision) (2.0.2)
Requirement already satisfied: pillow!=8.3.*,>=5.3.0 in /usr/local/lib/python3.11/dist-packages (from torchvision) (11.3.0)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.11/dist-packages (from jinja2->torch) (3.0.2)
2.6.0+cu124
CUDA is not available.
```

## Q6. How do you create a simple neural network in PyTorch?

```

import torch
import torch.nn as nn

# Define a simple neural network by subclassing nn.Module
class SimpleNN(nn.Module):
    def __init__(self, input_size, hidden_size, num_classes):
        super(SimpleNN, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size) # First fully connected layer
        self.relu = nn.ReLU() # ReLU activation function
        self.fc2 = nn.Linear(hidden_size, num_classes) # Second fully connected layer

    def forward(self, x):
        out = self.fc1(x)
        out = self.relu(out)
        out = self.fc2(out)
        return out

# Example Usage:
input_size = 10
hidden_size = 5
num_classes = 2

# Instantiate the model
model = SimpleNN(input_size, hidden_size, num_classes)

print("Model Architecture:")
print(model)

# Create a dummy input tensor
dummy_input = torch.randn(1, input_size) # Batch size of 1

# Perform a forward pass
output = model(dummy_input)

print("\nDummy Input:")
print(dummy_input)

print("\nOutput of the model:")
print(output)

```

```

🔍 Model Architecture:
SimpleNN(
  (fc1): Linear(in_features=10, out_features=5, bias=True)
  (relu): ReLU()
  (fc2): Linear(in_features=5, out_features=2, bias=True)
)

Dummy Input:
tensor([[ -0.4058, -0.2654,  0.2209,  0.0966, -0.3265,  0.9536, -1.1011, -0.7225,
          0.0225, -0.0225]])

Output of the model:
tensor([[0.1773, 0.2178]], grad_fn=<AddmmBackward0>)

```

## ✓ Q7. How do you define a loss function and optimizer in PyTorch?

```

import torch
import torch.nn as nn
import torch.optim as optim

# Assuming you have defined a model (like the SimpleNN from the previous example)
# For demonstration purposes, let's redefine a simple model here
class SimpleNN(nn.Module):
    def __init__(self, input_size, hidden_size, num_classes):
        super(SimpleNN, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size)
        self.relu = nn.ReLU()
        self.fc2 = nn.Linear(hidden_size, num_classes)

    def forward(self, x):
        out = self.fc1(x)
        out = self.relu(out)
        out = self.fc2(out)
        return out

input_size = 10
hidden_size = 5
num_classes = 2
model = SimpleNN(input_size, hidden_size, num_classes)

```

```
# 1. Define a Loss Function
# Common loss functions are available in the torch.nn module.
# For a classification task, Cross-Entropy Loss is often used.
criterion = nn.CrossEntropyLoss()

# For a regression task, Mean Squared Error (MSE) could be used:
# criterion = nn.MSELoss()

# 2. Define an Optimizer
# Optimizers are available in the torch.optim module.
# You need to pass the model's parameters to the optimizer.
# Stochastic Gradient Descent (SGD) is a basic optimizer.
# optimizer = optim.SGD(model.parameters(), lr=0.01)

# Adam is another popular optimizer.
optimizer = optim.Adam(model.parameters(), lr=0.001)

print("Loss Function:")
print(criterion)

print("\nOptimizer:")
print(optimizer)

# In a training loop, you would typically do the following steps in each iteration:
# 1. Forward pass: outputs = model(inputs)
# 2. Calculate loss: loss = criterion(outputs, labels)
# 3. Backward pass (compute gradients): loss.backward()
# 4. Update weights: optimizer.step()
# 5. Zero gradients: optimizer.zero_grad()
```

➡ Loss Function:  
CrossEntropyLoss()

Optimizer:  
Adam (  
Parameter Group 0  
amsgrad: False  
betas: (0.9, 0.999)  
capturable: False  
differentiable: False  
eps: 1e-08  
foreach: None  
fused: None  
lr: 0.001  
maximize: False  
weight\_decay: 0  
)

## ✓ Q8. How do you implement a custom loss function in PyTorch?

```
import torch
import torch.nn as nn
import torch.optim as optim

# Option 1: Implement a custom loss function as a Python function
def mean_squared_error_loss(predictions, targets):
    """Calculates the Mean Squared Error between predictions and targets."""
    return torch.mean((predictions - targets)**2)

# Option 2: Implement a custom loss function as a class inheriting from nn.Module
class MeanSquaredErrorLoss(nn.Module):
    def __init__(self):
        super(MeanSquaredErrorLoss, self).__init__()

    def forward(self, predictions, targets):
        """Calculates the Mean Squared Error between predictions and targets."""
        return torch.mean((predictions - targets)**2)

# Example Usage:
# Create some dummy predictions and targets
predictions = torch.randn(5, 1, requires_grad=True)
targets = torch.randn(5, 1)

# Using the function-based custom loss
loss_fn = mean_squared_error_loss(predictions, targets)
print(f"Loss using function: {loss_fn.item()}")

# Using the class-based custom loss
loss_module = MeanSquaredErrorLoss()
loss_class = loss_module(predictions, targets)
print(f"Loss using class: {loss_class.item()}")
```

```
# You can then use this custom loss in your training loop
# For example, for backpropagation:
# loss_class.backward()
# print(f"\nGradients of predictions after backward pass:\n{predictions.grad}")
```

```
↔ Loss using function: 1.237074613571167
   Loss using class: 1.237074613571167
```

## ✓ Q9. How do you save and load a TensorFlow model?

```
import tensorflow as tf
from tensorflow.keras import layers, models
import numpy as np
import os

# Create a simple sequential model (same as before)
model = models.Sequential([
    layers.Dense(units=16, activation='relu', input_shape=(10,)),
    layers.Dense(units=8, activation='relu'),
    layers.Dense(units=3, activation='softmax')
])

# Compile the model
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

# Generate some dummy data for demonstration
num_samples = 1000
input_dim = 10
num_classes = 3

x_train = np.random.rand(num_samples, input_dim).astype(np.float32)
y_train = np.random.randint(0, num_classes, num_samples)

# Train the model (briefly for demonstration)
print("Training the model...")
model.fit(x_train, y_train, epochs=10, verbose=0)
print("Model training finished.")

# Define a path to save the model with the recommended .keras extension
save_path = './my_tensorflow_model.keras'

# Save the entire model (architecture, weights, and optimizer state)
print(f"\nSaving the model to: {save_path}")
model.save(save_path)
print("Model saved successfully.")

# Load the model back
print(f"\nLoading the model from: {save_path}")
loaded_model = tf.keras.models.load_model(save_path)
print("Model loaded successfully.")

# You can now use the loaded model for predictions or further training
print("\nSummary of the loaded model:")
loaded_model.summary()

# Demonstrate prediction with the loaded model
dummy_input_for_prediction = np.random.rand(1, input_dim).astype(np.float32)
predictions = loaded_model.predict(dummy_input_for_prediction)
print(f"\nPrediction using loaded model for dummy input:\n{predictions}")

# Clean up the saved model file (optional)
# if os.path.exists(save_path):
#     os.remove(save_path)
#     print(f"\nRemoved saved model file: {save_path}")
```



```
Training the model...  
Model training finished.
```

```
Saving the model to: ./my_tensorflow_model.keras  
Model saved successfully.
```